

CSE-381: Systems 2

Exercise #4

Max Points: 20

You should save/rename this document using the naming convention **MUId_Exercise4.docx** (example: raodm_Exercise4.docx).

Objective: The objective of this exercise is to:

- Recap process states
- Use anonymous pipes to create software systems by interconnecting processes
- Named pipes aka FIFOs (First-In-First-Out).

Submission: Once you have completed this exercise, upload:

- This document saved as PDF file and named with the convention MUId_Exercise4.pdf.

You may discuss the questions with your instructor.

Name: Noah Dunn

Preliminaries



First wait for your instructor to cover the basic concepts of Inter-Process Communication and using anonymous pipes in Linux. Once your instructor has covered the concepts then you can proceed with this lab exercise.

The various command associated with this exercise are to be executed on the Linux server for this course. **You will not need to use NetBeans for this exercise.**

Part #1: Reviewing process states

Estimated time to complete: 10 minutes

1. Using the process state transition diagram discussed in class as a reference, indicate an example scenario when a process may undergo each one of the following state transitions during its life cycle:

State Transition	Example scenario or reason why a process may undergo such a state transition.
------------------	---

Running → Ready	When a program running on the CPU is interrupted during processing, it will go from running to ready state.
Running → Blocked	If the CPU is waiting on Input or Output to process or for an additional event to occur, CPU will swap from running to blocked
Blocked → Ready	When a required Input or Output is finished processing, or a required event finishes, CPU will swap from blocked to ready.
Ready → Running	At the discretion of the CPU scheduler, it will dispatch processes from ready to running as designed.

2. Compare and contrast preemptive vs. non-preemptive multitasking. In separate sentences, state one advantage and one disadvantage of each of the two approaches.

Non-preemptive Multitasking:

Advantage- OS is easier to develop

Disadvantage- Malicious applications can take over CPU

Preemptive Multitasking:

Advantage- Fault tolerant to malicious or buggy applications

Disadvantage- OS is harder to develop

Part #2: Using anonymous pipes (aka FIFOs) [5 points]

Estimated time to complete: 25 minutes

Background: Anonymous pipes provide a convenience mechanism for IPC between processes running on the same computer. Anonymous pipes are also convenient mechanisms to interconnect processes on a Linux machine to perform tasks -- **giving raise to a software system or a "software pipeline"**. Most of the command-line programs and utilities in Linux are designed to facilitate the use of anonymous pipes as they read and write data to standard input (`std::cin`) and standard output (`std::cout`) respectively. The `bash` shell enables "wiring" or "connecting" standard I/O streams via the `|` (pipe) operator to setup utility software pipelines for performing sophisticated tasks as shown in the example below:

List all processes for user raodm

```
$ ps -fe | grep raodm
```

There are too many useful Linux commands/tools to cover, but here are a few common ones that you will be using in this exercise:

Command	Description	Example
cut	Cut out columns specified by -f ; columns delimited by 1 character specified by -d	\$ cut -d "," -f 1,2,4 faculty.txt \$ ps -fea cut -d " " -f1
grep	Search for string/regular-expression in file or standard input	\$ grep "ra?dm" test.txt
sort	Sort file/inputs based on key column specified via -k ; columns delimited by 1 character specified via -t	\$ sort -k 2 -t , faculty.txt \$ cut -d" " -f2,3 classes.txt sort -t, -k 1
join	Join 2 files based on common values in 2 given columns. Each file must be sorted on common value/column specified by -1 (for file 1) and -2 (for file 2); columns delimited by 1 character specified via -t	\$ join -t, -1 2 -2 1 classes.txt courses.txt \$ sort -k2 -t, classes.txt join -t, -1 2 -2 1 - courses.txt

Exercise: This exercise is purely a problem solving one that requires you to innovate by using the above 4 commands to perform the following operations (that would typically be performed in an interview setting as well):

1. For this exercise you are supplied with the following 3 Comma Separated Values (CSE) files: `classes.txt`, `courses.txt`, and `faculty.txt`. These are simple files. View the files to understand the data in them. Scp these 3 files to the Linux server.

Next collaboratively develop a sequence of 1 or more commands (commands are connected via pipes) to extract the data for the following problems. Sample outputs are show for each problem below.

2. List just the faculty names (Hint: `cut`):

as

Sample output:

```
Campbell  
Kiper  
Rao
```

Copy-paste **command and the output** from your command in space below:

```
dunnm2@os1:~$ cut -c 5- faculty.txt  
Campbell  
Kiper  
Rao
```

3. List the classes (as in cse278, cs301, etc.) taught in Fall17 (Hint: grep + cut)

Sample output:

```
cse278
cse278
cse301
cse443
```

Copy-paste **command and the output** from your command in space below:

```
dunnnm2@os1:~$ grep Fall17,cse classes.txt | cut -c 8-13
cse278
cse278
cse301
cse443
```

4. List the classes and names of courses taught in Spring18:

Sample output:

```
cse278,Systems 1
cse301,Software Architecture
cse381,Systems 2
```

Note: The following command will come in handy. Run the command below and observe how the join command combines 2 files based on the values of the course number (e.g., cse278)

```
join -t, -1 2 -2 1 classes.txt courses.txt
```

Copy-paste command and the output from your command in space below:

```
dunnnm2@os1:~$ grep Spring18 classes.txt | join -t, -1 2 -2 1 -
courses.txt | cut -d , -f 1,4
cse278,Systems 1
cse301,Software Architecture
cse381,Systems 2
```

5. Optional challenge (come back to this one at the end if you have time): List classes, course name, and faculty names of courses taught in Spring18:

Sample output:

```
cse278,Systems 1,Campbell
cse301,Software Architecture,Kiper
cse381,Systems 2,Rao
```

Copy-paste command and the output from your command in space below:

```
dunnnm2@os1:~$ grep Spring18 classes.txt | join -t, -1 2 -2 1 -  
courses.txt | join -t, -1 3 -2 1 - faculty.txt | cut -d , -f  
2,4,5  
cse278,Systems 1,Campbell  
cse301,Software Architecture,Kiper  
cse381,Systems 2,Rao
```

Part #3: Using named pipes (aka FIFOs) [5 points]

Estimated time to complete: 25 minutes

Background: A named pipe (also known as a FIFO because it performs First-In-First-Out type IPC) is a commonly used pipe concept on Linux and Unix-like systems. It is one of the different methods for inter-process communication. The concept of named pipes is also found in Microsoft Windows, although the semantics differ substantially. Unlike a traditional pipe that is "unnamed" (as it exists anonymously and persists only for as long as the interacting process is running, named pipes exist beyond the lifetime of processes), a named pipe is system-persistent and exists beyond the life of the process. It can be deleted once it is no longer being used. Processes generally attach to the named pipes **as if they were normal files** to perform inter-process communication (IPC) by reading and writing data to it. Usually pipes are used to exchange text data. However, the pipes may also be used for interchange of binary data.

Exercise: This exercise requires you to create a FIFO and use it for exchanging data between two processes via the following experimental procedure:

1. Create a FIFO using the following command at the shell prompt (where MUIId is your Miami University unique ID):

```
$ mkfifo /tmp/MUIId
```

Next, setup the permission on the pipe using `chmod` such that: user (i.e., you) can read-and-write while group and other users can only write to your FIFO (i.e., `/tmp/MUIId`). To verify the named pipe has been set with correct permissions, run `ls -l /tmp/MUIId` to list the file entry and verify the permissions.

Copy-paste the `chmod` command and use the `ls -l` output to validate the permissions on the pipe have been set correctly:

```
$ chmod u+rw-x,og+r-wx /tmp/raodm  
$ ls -l /tmp/raodm  
prw-r--r-- 1 raodm uidd 0 Feb 21 13:03 /tmp/raodm
```

2. Now use the FIFO that was just created to perform some simple IPC between two different processes. For this we will create one process that sends data to the FIFO and another one that reads data from the FIFO as directed below:

- a. Open two different terminal window and log into the Linux server (now you will have 2 different terminals, both logged onto the Linux server).
 - b. In one of the windows print the contents of the FIFO using the following command (where *MUId* is the FIFO that was created earlier):

```
$ cat /tmp/MUId
```
 - c. In the other terminal write data to the FIFO using the following command:

```
$ ps -fe > /tmp/MUId
```
 - d. You should observe the output displayed in window in which the earlier `cat` command was run.
3. Observe how the FIFO enables communication between processes using following steps:
- a. Run “`ls -l`” and record the size (**number before date**) of the FIFO in the space below:

Size (in bytes) of the FIFO was: 0
 - b. Don't run `cat` to read data of FIFO. Instead, just write output of `ps -fe` command to the named pipe (as shown earlier). This process will appear to hang, which is normal -- this is because the FIFO's internal memory buffer is full and it is waiting for another process to read data to free-up space.
 - c. While the `ps -fe` command is still waiting to complete, from another terminal to check the size of the FIFO reported by “`ls -l`”. Its size should remain unchanged (because all the data is stored in memory by the OS).
 - d. Finally, you can `cat` the data in your FIFO and observe the full output `ps` is returned as the data was stored in internal OS buffers.

Play with permissions. Now remove write permissions on your FIFO for everyone using the `chmod` command. Verify permission settings on the FIFO using `ls -l` command:

Copy-paste the `chmod` command and use the `ls -l` output to validate the permissions on the pipe have been set correctly:

```
$ chmod a-w /tmp/raodm
$ ls -l /tmp/raodm
prw-r--r-- 1 raodm uidd 0 Feb 21 13:03 /tmp/raodm
```

4. Try to write some data to the FIFO and ensure that permission is denied to write data to the FIFO.
5. You may re-enable write permissions on the FIFO for everyone and try using the named pipe.

6. Ask your neighbor to write some data to your named pipe and see how different users can write data to it.
7. Write a short message to your instructors FIFO.
8. Finally delete the FIFO that you created using the `rm` command as shown below:

```
$ rm /tmp/MUId
```

Part #4: Submit files to Canvas

Save this document (duly filled with the necessary information) as a PDF file using the naming convention `MUId_Exercise4.pdf`. Upload the PDF file to Canvas.

Upload each file individually onto Canvas. Do not upload `zip/7zip/tar/gz` or other archive file formats for your submission.